

In lieu of cover sheet:

Name: Peter Lotts

CRSid: pml43

College: Downing

Director of Studies: Dr Robert Harle

Overseers: Prof. J. Daugman & Dr. D J Greaves

Title of project: Adding network subsystem provenance collection to
CADETS

Originator: Dr R. Sohan

Supervisor: Dr G. Jenkinson

Introduction and Description of Work

Nowadays, most cyberattacks from nation state actors do not take on the traditional form of a piece of malware which tries to find a store of sensitive data and dump it back to the attacker as quickly as possible before system administrators have time to do anything. Instead, these attacks slowly infiltrate the target system, hiding out of sight and learning the use patterns of normal users. It then exploits this knowledge to slowly make horizontal transfers within the target enterprise's internal network, and extract sensitive data from databases in such a way that the access is hard to distinguish from normal access. Such an attack is called an Advanced Persistent Threat (APT), and can go on for months or years without system administrators realising. Eventually, it is likely that the malware will make a mistake and trigger an indicator of compromise to be observed by administrators, but by then it is difficult to tell when and how the malware entered the system, and what data it has seen. This can make impact assessments, as well as efforts to prevent a similar attack in future, very difficult.

The Computer Laboratory has a project which is trying to combat this by building CADETS, a Causal, Adaptive, Distributed, and Efficient Tracing System. The system is being built to track the provenance of data by collecting metadata from computers all over the network about what processes act on what data and when. This metadata is then collected in a distributed database, where it can be analysed to trace data flows throughout the computer system. This should allow analysts to trace malware back to possible entry times and methods, which they can then analyse further. This project will add support for collecting metadata on network packets as they flow through the kernel network stack.

Starting Point

The CADETS project currently has a user-space application which collects metadata on kernel-level datastructures via libdtrace, translates the metadata to a json format for easy interpretation, and then sends this away for processing (often over the network). It also has an application with a user interface to display the data it has collected.

The FreeBSD kernel provides a means of tagging its main internal structure of interest, namely *struct mbuf* using *mbuf_tags*. It is thought that this will be sufficient to tag a packet's data with a unique identifier in order to track its progress through the kernel's network stack.

Under Linux, the *Resourceful* framework is able to collect data from auto-generated tracepoints within the Linux kernel with relatively low overhead, but it does not currently inspect the network subsystem.

Substance and Structure of the Project

The objective of this project is to make additions to the DTrace tracing system under FreeBSD to allow it to trace the flow of data through the network subsystem of the FreeBSD kernel, and allow this data to be integrated with the wider CADETS project. The project should be able to track kernel memory allocations which occur as a result of network packet flow, along with times taken for the different network layers to be traversed and attempting to estimate the cause of delays, such as scheduling conflicts. This extension will allow DTrace probes to access the new packet metadata, and it will therefore be accessible to the existing CADETS userspace application which gathers data for processing.

CADETS already generates a locally unique identifier for each socket opened by a process under observation. In order to track packets through the network stack, this project will need to tag each packet with a further unique identifier, and maintain an association between this identifier and that of the related socket. The FreeBSD kernel already provides a method for tagging kernel structures called *mbufs*, which hold the data payload of a packet as it flows through the network layers. These can be used as the mechanism for attaching an identifier to each packet. The significant complexity in this area of the project might come from the generation of locally unique identifiers to form tags, as this process must be carried out in a timely fashion in order to prevent unacceptable delays in the processing of network packets. The FreeBSD kernel provides methods for generating unique identifiers, but their performance will have to be assessed to determine whether they are fast enough to be of use. If they are not, then an alternative system, based on existing open source algorithms, will need to be developed.

Metadata collection will be performed by extending DTrace's existing collection schemes, aiming to modify this code rather than writing native kernel code where possible, as this will increase portability across new kernel versions. Once this data is made available to DTrace, the CADETS userspace code should be able to collect it for further processing. This userspace code may need to be modified a little, as it may be most useful for events pertaining to each packet to be reconciled together and presented as a summary record, whereas the incoming data from DTrace will be split according to which CPU core the kernel executed the tracepoint on.

To evaluate the performance impact that the project's trace points are having on the performance of the kernel's network subsystem, the network interface throughput will be measured at the socket level without any monitoring, then with DTrace enabled and collecting general network information but without this project's additions, and then with DTrace collecting data from this project's additional code. This will allow the performance drop to be estimated in context with the performance impact of running DTrace.

Extensions

If the core part of the project is completed in sufficient time, there are several extensions which could be implemented:

- 1) A simple visualisation tool could be produced as a DTrace consumer in place of the current CADETS userspace application. This would allow the results of monitoring to be viewed in a user-friendly format.
- 2) Based on trace output from using the full tracing system, the most commonly executed pieces of new code could be found and optimised, if they are deemed to be adversely affecting performance. It may be necessary to use some of the concepts from the *Resourceful* project if DTrace itself is deemed to be causing significant overhead. Any *Resourceful* code will need to be ported from Linux to FreeBSD if it is to be used directly.
- 3) The existing CADETS user interface could be extended to allow the newly collected packet data to be viewed, and to allow users to dig into the flow of data from a particular socket which is under investigation.

Success Criterion

The project will be successful if it is able to collect metadata on the data flows within the FreeBSD kernel's network stack, associating the journey each packet makes with the socket to which it is linked. This should be achieved without incurring an unacceptable performance overhead for the delivery of network packets.

Plan of Work

The main bulk of the project (after the proposal has been formally accepted) will be split into 14 2-week work packages, with a final week of contingency time at the end of the project before the dissertation submission deadline. Note that the work package over the Christmas period is extended to 3 weeks to allow for inevitable time away from the project during this time. I am planning to take lecture courses with a fairly even split between the Michaelmas and Lent terms, so the workload should be quite even throughout the project. For project management purposes, the expected milestones from each work package will be entered into Bitbucket's issue tracker, allowing progress to be estimated at every stage of the project.

Before 20/10/2017

Find project supervisor and discuss ideas, submit Phase 1 Report form to Overseers and gain their initial approval. Make contact with the Computer Laboratory group developing the CADETS project and meet to discuss ideas. Further discussion with Overseers regarding details of the project. Submit draft proposal, and continue to refine this based on comments from Overseers. Submit final proposal by 12:00 20/10/2017.

21/10/2017 – 03/11/2017

Following acceptance of the project proposal, the research phase of the project can commence. This will involve studying how DTrace is used, and how it can be extended to add more trace points yielding more information. The expected output from this process is some small DTrace scripts, written in DTrace's D language, to collect data about sockets by the methods currently provided by DTrace.

04/11/2017 – 17/11/2017

With some familiarity of DTrace and what is possible, the implementation of the FreeBSD kernel can now be inspected to identify where the kernel should have tracepoints added, what information these would make available and whether this is sufficient, and where each packet's *mbuf* can have its identifier generated and be tagged. The expected output of this process is a list of references to kernel source lines which look to be good candidates for adding instrumentation.

18/11/2017 – 01/12/2017

The remaining research is into the UUID generation algorithms available in the kernel. Studying these may require manual benchmarking code to be written to assess their performance, and therefore whether they would be suitable for tagging packets. If they are determined to be too slow, then other algorithms will need to be researched and one selected and implemented. Once an algorithm has been chosen, this can be used to implement packet tagging in the kernel.

02/12/2017 – 15/12/2017

Leave Cambridge. Work on adding tracepoints at the locations previously identified, making use of the tags which packets now have attached to them in the kernel. At the top of the kernel stack, make sure that packet identifiers can be associated with their socket.

16/12/2017 – 05/01/2018

Set up useful DTrace probes using new tracepoints to collect data for import into CADETS. This is the phase where any required modifications in the userspace CADETS code will be made to rationalise the incoming data from the different CPU core streams.

06/01/2018 – 19/01/2018

Return to Cambridge. Begin writing progress report and run throughput benchmarking in the three different configurations in order to make a performance assessment of the project code.

20/01/2018 – 02/02/2018

Use results from benchmarking to inform decisions about any project extensions which may be completed, and whether further code optimisation is needed. Complete progress report and send to Supervisor and Overseers. Progress report submission deadline: 12:00 02/02/2018

03/02/2018 – 16/02/2018

Complete functionality of core project. Perform any major refactoring which is deemed necessary, and begin to look at the possible extensions to the project.

17/02/2018 – 02/03/2018

Contingency time for core project / time to work on project extensions. Begin collecting ideas about dissertation including code samples to be used.

03/03/2018 – 16/03/2018

Contingency time for core project / time to work on project extensions.

17/03/2018 – 30/03/2018

Leave Cambridge. Code should by now be functional and approaching its final form. Work on dissertation commences in earnest. Write Introduction section, and begin writing framework for the Preparation and Implementation sections. Code completion deadline: 30/03/2018.

31/03/2018 – 13/04/2018

Complete Preparation and Implementation sections of dissertation from existing frameworks. Finish collecting data for the Evaluation section if need be. Code should be complete, barring small aesthetic changes found to be necessary when writing the Implementation section of the dissertation.

14/04/2018 – 27/04/2018

Return to Cambridge. Complete Evaluation and Conclusions sections of dissertation, meeting with supervisor in person to discuss the current draft of the whole dissertation. Start to check over the required sections around the dissertation itself (table of contents, bibliography, etc).

28/04/2018 – 11/04/2018

Make final changes to dissertation with project supervisor, and ask Director of Studies to read and comment on it. From this, the final copy of the dissertation may be produced.

12/04/2018 – 18/04/2018

Contingency time in case the dissertation is not yet fully complete. This time may be filled with final alterations or more significant writing if this is required. Deadline for dissertation submission: 12:00 18/04/2018.

Resources Declaration

The languages used for this project are entirely determined by the existing code with which the project will be integrating. The kernel module component will have to be written in C, to match the FreeBSD kernel, and DTrace scripts must be written in the D scripting language. This project does not depend on specific hardware, but will require an installation of FreeBSD to test on. I will use a virtual machine for this as I do not have easy access to a computer natively running FreeBSD, and it is easy to take regular backups of a virtual machine's hard drive.

For ease of access, I intend to use my own computer for most development and as the host machine for my virtual testing. It's specifications are: Intel i5-5200U Processor (2.2GHz dual core with hyperthreading), 16GB RAM, Windows 10, VirtualBox 5.1.28 at time of writing (updates will be applied throughout the project as they become available). I accept full responsibility for this machine and I have made contingency plans to protect myself against hardware and/or software failure. Once the virtual machine is set up, a full backup of the host machine will be taken to external media, allowing its state to be restored to the project start state.

All code will be subject to version control using *git*, and the repository will be pushed to the Bitbucket cloud service regularly. Periodically, a copy will also be taken onto the MCS file space, ensuring that a copy is always available in Cambridge. Any files not committed to git will be backed up regularly to external media.